

**Smeet Sheth**  
**Data 225 HW5**  
**016786133**

### Getting started with Snowflake:

The screenshot displays the Snowflake user interface. On the left, a user menu is open for 'Smeet Sheth' (ACCOUNTADMIN). The menu includes options for 'Switch Role' (ACCOUNTADMIN), 'Profile', 'Documentation', 'Support', 'Privacy Notice', and 'Sign Out'. On the right, the 'Worksheets' section is visible, showing a table with columns 'TITLE' and 'TYPE'. The table lists several roles: ACCOUNTADMIN (selected), ORGADMIN, PUBLIC, SECURITYADMIN, SYSADMIN, and USERADMIN.

TITLE	TYPE
ACCOUNTADMIN	✓
ORGADMIN	
PUBLIC	
SECURITYADMIN	
SYSADMIN	
USERADMIN	

Query for creating the hotel\_bookings table:

### View Definition

```
1  create or replace TABLE DBHW5.PUBLIC."hotel_bookings" (  
2      HOTEL VARCHAR(100),  
3      IS_CANCELED NUMBER(38,0),  
4      LEAD_TIME NUMBER(38,0),  
5      ARRIVAL_DATE_YEAR NUMBER(38,0),  
6      ARRIVAL_DATE_MONTH VARCHAR(20),  
7      ARRIVAL_DATE_WEEK_NUMBER NUMBER(38,0),  
8      ARRIVAL_DATE_DAY_OF_MONTH NUMBER(38,0),  
9      STAYS_IN_WEEKEND_NIGHTS NUMBER(38,0),  
10     STAYS_IN_WEEK_NIGHTS NUMBER(38,0),  
11     ADULTS NUMBER(38,0),  
12     CHILDREN NUMBER(38,0),  
13     BABIES NUMBER(38,0),  
14     MEAL VARCHAR(100),  
15     COUNTRY VARCHAR(100),  
16     MARKET_SEGMENT VARCHAR(100),  
17     DISTRIBUTION_CHANNEL VARCHAR(100),  
18     IS_REPEATED_GUEST NUMBER(38,0),  
19     PREVIOUS_CANCELLATIONS NUMBER(38,0),  
20     PREVIOUS_BOOKINGS_NOT_CANCELED NUMBER(38,0),  
21     RESERVED_ROOM_TYPE VARCHAR(100),  
22     ASSIGNED_ROOM_TYPE VARCHAR(100),  
23     BOOKING_CHANGES NUMBER(38,0),  
24     DEPOSIT_TYPE VARCHAR(100),  
25     AGENT VARCHAR(100),  
26     COMPANY VARCHAR(100),  
27     DAYS_IN_WAITING_LIST NUMBER(38,0),  
28     CUSTOMER_TYPE VARCHAR(100),  
29     ADR FLOAT,  
30     REQUIRED_CAR_PARKING_SPACES NUMBER(38,0),  
31     TOTAL_OF_SPECIAL_REQUESTS NUMBER(38,0),  
32     RESERVATION_STATUS VARCHAR(100),  
33     RESERVATION_STATUS_DATE DATE  
34 );
```

Imported the data into our table:

### Data Partially Loaded into Table

hotel\_bookings.csv → hotel\_bookings

119386 of 119,391 rows were successfully inserted into the table.

The first error occurs on line 1 at character 7 on column 2 ("IS\_CANCELED").

Numeric value 'is\_canceled' is not recognized

[See more and explore errors in Worksheets](#)

Query Data

Done

Created a Warehouse for the project:

< DB225WAREHOUSE ...

Warehouse

ACCOUNTADMIN

36 minutes ago

#### Details

Status	Size	Max Clusters
Suspended	X-Small	1
Min Clusters	Scaling Policy	Running
1	STANDARD	0
Queued	Auto Suspend	Auto Resume
0	600 seconds	Enabled
Resumed On	Query Acceleration	Scale Factor
18 minutes ago	Disabled	8

Query History

Status All

User All

Filters

52 Queries

Columns

SQL TEXT	QUERY ID	STATUS	USER
with active_contracts as ( select disti...	01ac10e4-0001-9d9a-0022-2307000130be	Success	SMEET42
with active_contracts as ( select disti...	01ac10e1-0001-9d3a-0022-2307000110ca	Success	SMEET42
SELECT TOP 1 QUERY_ID FROM TABLE(DBHW5...	01ac10de-0001-9da7-0022-2307000100aa	Success	SMEET42
with active_contracts as ( select disti...	01ac10c1-0001-9d3a-0022-23070001102e	Success	SMEET42
SELECT * FROM "hotel_bookings" limit 5;	01ac10b8-0001-9d71-0000-002223075279	Success	SMEET42
with active_contracts as ( select disti...	01ac10ac-0001-9d71-0000-002223075275	Success	SMEET42
SELECT hotel, customer_type, COUNT(*) A...	01ac0e0a-0001-9d7b-0000-0022230732...	Success	SMEET42
SELECT CASE WHEN lead_time <= 7 THEN '0...	01ac0e08-0001-9d7b-0000-0022230732...	Success	SMEET42

Created Internal stage to extract data from the database:

Databases Worksheets

ACCOUNTADMIN DB225WAREHOUSE Share

Pinned (1)  
Supply\_Chain

DBHW5  
INFORMATION\_SCHEMA  
PUBLIC  
Tables  
Stages  
HOTELWAREHOUSE  
SNOWFLAKE  
SNOWFLAKE\_SAMPLE\_DATA

DBHW5.PUBLIC Settings

1 create stage hotelwarehouse;

Results Chart

status
1 Stage area HOTELWAREHOUSE successfully created.

Query Details

Query duration 213ms  
Rows 1  
Query ID 01ac10c4-0001-9d9a-0...

Displaying the stage we created and its details:

DBHW5.PUBLIC ▾

Settings ▾

LATEST VERSION ▾ 🔍

1

`show stages;`

Results Chart 🔍 📄 ⬇️ 📄

	...	created_on	name	database_name	schema_name	url	has_credentials	has_encrypt
1		2023-05-04 10:40:50.152 -0700	HOTELWAREHOUSE	DBHW5	PUBLIC		N	N

Stage details:

SS Smeet Sheth  
ACCOUNTADMIN ▾

Worksheets

Dashboards

Data

**Databases**

Private Sharing

Provider Studio

Marketplace

Activity

Admin

Get Started

Help & Support

⌚

29 days left in trial

🔍 Search

...

DBHW5

INFORMATION\_SCHEMA

PUBLIC

Tables

Stages

**HOTELWAREHOUSE**

SNOWFLAKE

SNOWFLAKE\_SAMPLE\_DATA

📁 DBHW5 / PUBLIC / HOTELWAREHOUSE

⋮ + Files

📁 Internal Stage

👤 ACCOUNTADMIN

🕒 4 minutes ago

Stage Files

**Stage Details**

Details

Storage Integration  
Internal

Privileges

Group by Role ▾ + Privilege

👤 ACCOUNTADMIN (Current Role)

🔍 OWNERSHIP

## Creating file format that will be uploaded into the warehouse:

The screenshot shows the Snowflake web interface. On the left, the 'Databases' tab is active, and the object browser shows the hierarchy: DBHW5 > PUBLIC > File Formats. The 'Hotel\_CSV\_Format' file format is listed. The main query editor shows the following SQL code:

```
1 create or replace file format "Hotel_CSV_Format"
2 type = 'CSV'
3 field_delimiter = ','
4
```

The 'Results' tab is selected, showing a single row with the status: 'File format Hotel\_CSV\_Format successfully created.' The 'Query Details' panel on the right shows: Query duration: 80ms, Rows: 1, and Query ID: 01ac10d0-0001-9d7b-0...

## Creating table 'hotelbookingdata' inside stage:

The screenshot shows the Snowflake web interface. On the left, the 'Databases' tab is active, and the object browser shows the hierarchy: DBHW5 > PUBLIC > Stages > HOTELWAREHOUSE. The 'hotel\_bookings' table is listed. The main query editor shows the following SQL code:

```
1 copy into @HOTELWAREHOUSE/hotelbookingdata
2 from "hotel_bookings"
3 file_format = (format_name = "Hotel_CSV_Format");
```

The 'Results' tab is selected, showing a table with the following data:

	rows_unloaded	input_bytes	output_bytes
1	119,386	16,854,490	1,215,351

Table created successfully:

DBHW5.PUBLIC ▾

Settings ▾

LATEST VERSION ▾ 🔍

1

`list @HOTELWAREHOUSE/hotelbookingdata`

↶ Results

Chart

🔍 📄 ⬇️ 🗑️

	name	size	md5	...	last_modified
1	hotelwarehouse/hotelbookingdata_0_0_0.csv.gz	1,215,360	4472eb04cea4d005f4d0df4e9335221f		Thu, 4 May 2023 17:57:55 GMT

New table created inside stage:

SS Smeet Sheth  
ACCOUNTADMIN ▾

Worksheets

Dashboards

Data

Databases

Private Sharing

Provider Studio

Marketplace

Activity

Admin

Get Started

Help & Support

Search

DBHW5

INFORMATION\_SCHEMA

PUBLIC

Tables

Stages

HOTELWAREHOUSE

File Formats

SNOWFLAKE

SNOWFLAKE\_SAMPLE\_DATA

DBHW5 / PUBLIC / HOTELWAREHOUSE

⋮ + Files

Internal Stage

ACCOUNTADMIN

26 minutes ago

Stage Files

Stage Details

HOTELWAREHOUSE (2 Files)

PREVIEW

🔍 Search

+ COMPUTE\_WH

🔄

NAME	SIZE	LAST MODIFIED	⬇️
hotelbookingdataold_0_0_0.csv.gz	1.2MB	1 minute ago	⋮
hotelbookingdata_0_0_0.csv.gz	1.2MB	9 minutes ago	⋮

Removing the new table:

DBHW5.PUBLIC ▾

Settings ▾

LATEST VERSION ▾ 🔍

1

`remove @HOTELWAREHOUSE/hotelbookingdataold`

2

↶ Results

Chart

🔍 📄 ⬇️ 🗑️

	name	...	result
1	hotelwarehouse/hotelbookingdataold_0_0_0.csv.gz		removed

## Queries for analysis:

### Query 1.

```
SELECT
  hotel,
  COUNT(*) AS total_bookings,
  SUM(is_canceled) AS total_cancellations,
  SUM(is_canceled) / COUNT(*) AS cancellation_rate
FROM
  "hotel_bookings"
GROUP BY
  hotel;
```

↶ Results

↗ Chart

	HOTEL	...	TOTAL_BOOKINGS	TOTAL_CANCELLATIONS	CANCELLATION_RATE
1	Resort Hotel		40,060	11,122	0.277634
2	City Hotel		79,326	33,098	0.41724

This query calculates the cancellation rate of two different categories of hotels.

### Query 2.

```
SELECT
  sub.hotel,
  sub.reserved_room_type,
  sub.assigned_room_type,
  sub.booking_status,
  sub.avg_daily_rate
FROM
  (
    SELECT
      hotel,
      reserved_room_type,
      assigned_room_type,
      CASE
        WHEN is_canceled = 0 THEN 'booked'
        ELSE 'canceled'
      END AS booking_status,
      AVG(adr) AS avg_daily_rate
```



```

FROM
    "hotel_bookings"
GROUP BY
    hotel,
    reserved_room_type,
    assigned_room_type,
    booking_status
) sub
WHERE
    sub.hotel IN (
        SELECT
            hotel
        FROM
            "hotel_bookings"
        WHERE
            is_canceled = 0
        GROUP BY
            hotel
        ORDER BY
            AVG(adr) DESC
        LIMIT 10
    )
ORDER BY
    sub.avg_daily_rate DESC

limit 5;

```

Results		Chart							
	HOTEL	RESERVED_ROOM_TYPE	ASSIGNED_ROOM_TYPE	BOOKING_STATUS	AVG_DAILY_RATE				
1	Resort Hotel	F	D	canceled	289.6				
2	City Hotel	G	F	canceled	273				
3	Resort Hotel	G	C	booked	268				
4	City Hotel	G	G	canceled	226.315384615				
5	Resort Hotel	D	C	canceled	208.67				

The above query calculates the average daily rate for all the hotels and displays the top 5 hotels with highest ADR along with their booking status. We can use the query to understand the functioning of these hotels and apply them to the hotels with comparatively low ADR

### Query 3.

```
SELECT
  h.hotel,
  AVG(h.lead_time) AS avg_lead_time,
  AVG(h.stays_in_week_nights + h.stays_in_weekend_nights) AS avg_stay_length
FROM
  "hotel_bookings" h
  INNER JOIN (
    SELECT
      hotel,
      arrival_date_month,
      AVG(adr) AS avg_monthly_adr
    FROM
      "hotel_bookings"
    WHERE
      is_canceled = 0
    GROUP BY
      hotel,
      arrival_date_month
  ) m ON h.hotel = m.hotel AND h.arrival_date_month = m.arrival_date_month
WHERE
  h.is_canceled = 0
  AND h.assigned_room_type = 'A'
  AND h.arrival_date_year = 2017
GROUP BY
  h.hotel;
```

Results			
	HOTEL	AVG_LEAD_TIME	AVG_STAY_LENGTH
1	Resort Hotel	82.672908	4.263546
2	City Hotel	102.593308	2.985768

The above query calculates the Average stay of the customers for both the type of hotels in the year 2017 with room type 'A'. This can help understand the customer pattern for that particular hotel and can be analyzed to provide services accordingly.

#### Query 4.

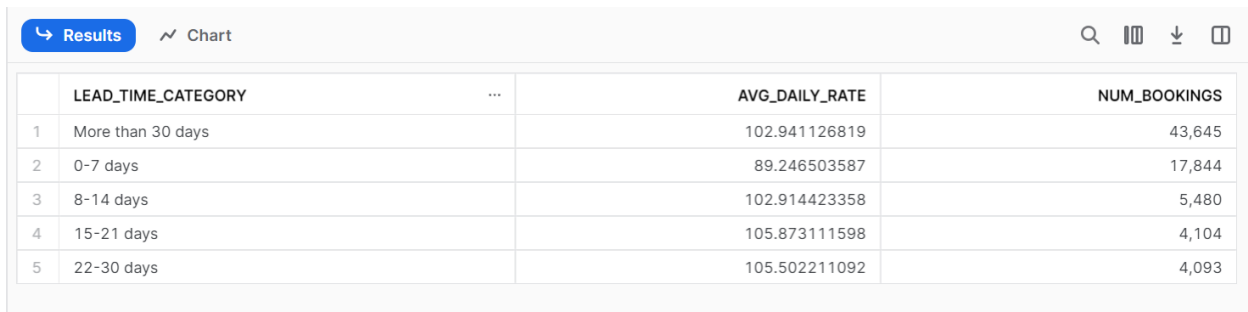
```
SELECT
    hotel,
    arrival_date_month,
    AVG(adr) AS avg_daily_rate,
    RANK() OVER (PARTITION BY arrival_date_month ORDER BY AVG(adr) DESC) AS
rank_by_month
FROM
    "hotel_bookings"
WHERE
    is_canceled = 0
GROUP BY
    hotel,
    arrival_date_month;
```

Results		Chart			
	HOTEL	ARRIVAL_DATE_MONTH	...	AVG_DAILY_RATE	RANK_BY_MONTH
1	City Hotel	September		112.598452214	1
2	Resort Hotel	September		96.416860133	2
3	Resort Hotel	August		181.205891925	1
4	City Hotel	August		118.412083256	2
5	City Hotel	November		86.500456231	1
6	Resort Hotel	November		48.681639676	2
7	City Hotel	December		87.856764214	1
8	Resort Hotel	December		68.322235994	2
9	City Hotel	January		82.160634428	1
10	Resort Hotel	January		48.70891863	2
11	City Hotel	February		86.183025457	1
12	Resort Hotel	February		54.147478336	2
13	Resort Hotel	July		150.122527893	1
14	City Hotel	July		115.563810121	2

The query uses a window function to compare the ADR for every month and rank them accordingly. Further analysis can be done to understand the pattern of the customers which are changing based on the change in seasons.

### Query 5.

```
SELECT
CASE
  WHEN lead_time <= 7 THEN '0-7 days'
  WHEN lead_time <= 14 THEN '8-14 days'
  WHEN lead_time <= 21 THEN '15-21 days'
  WHEN lead_time <= 30 THEN '22-30 days'
  ELSE 'More than 30 days'
END AS lead_time_category,
AVG(adr) AS avg_daily_rate,
COUNT(*) AS num_bookings
FROM
  "hotel_bookings"
WHERE
  is_canceled = 0
GROUP BY
  lead_time_category;
```



	LEAD_TIME_CATEGORY	AVG_DAILY_RATE	NUM_BOOKINGS
1	More than 30 days	102.941126819	43,645
2	0-7 days	89.246503587	17,844
3	8-14 days	102.914423358	5,480
4	15-21 days	105.873111598	4,104
5	22-30 days	105.502211092	4,093

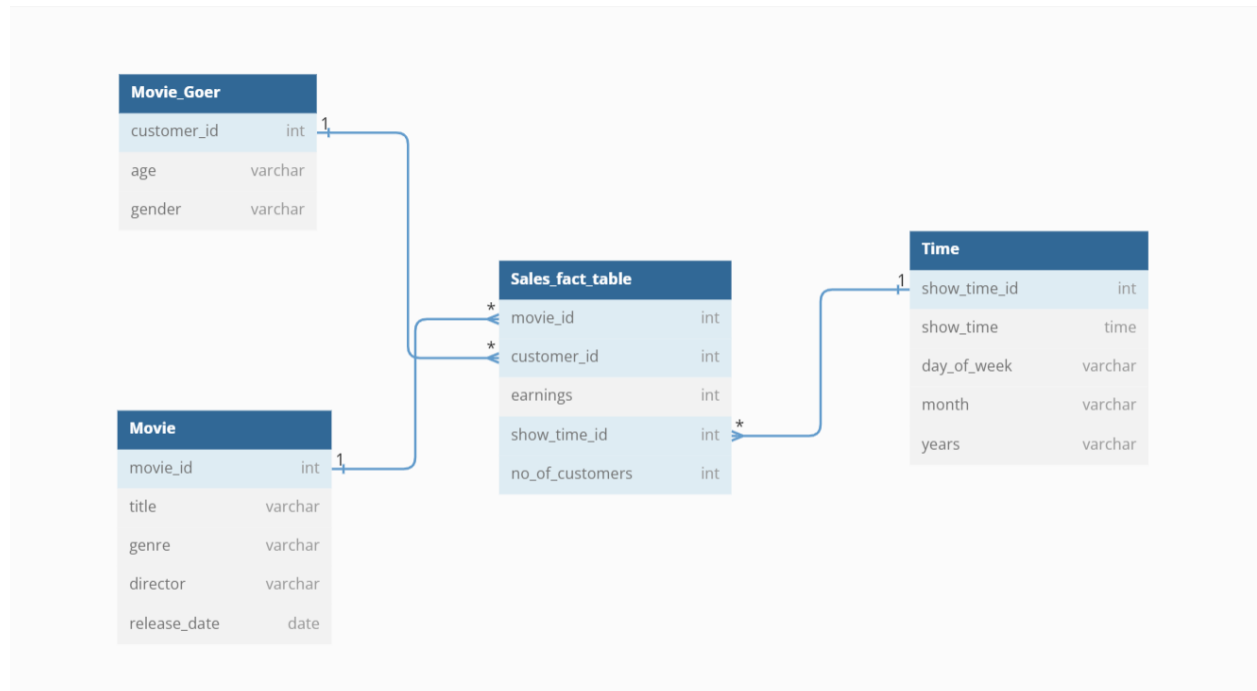
This query categorizes the ADR for different lengths of stay of the customers and also compare the total number of bookings that fall under each category.

### Q1 B:

#### Blog link:

<https://medium.com/@smeetsheth2001/data-warehousing-using-snowflakeuser-experience-fbf5b3a06a31>

**Q2 A:**



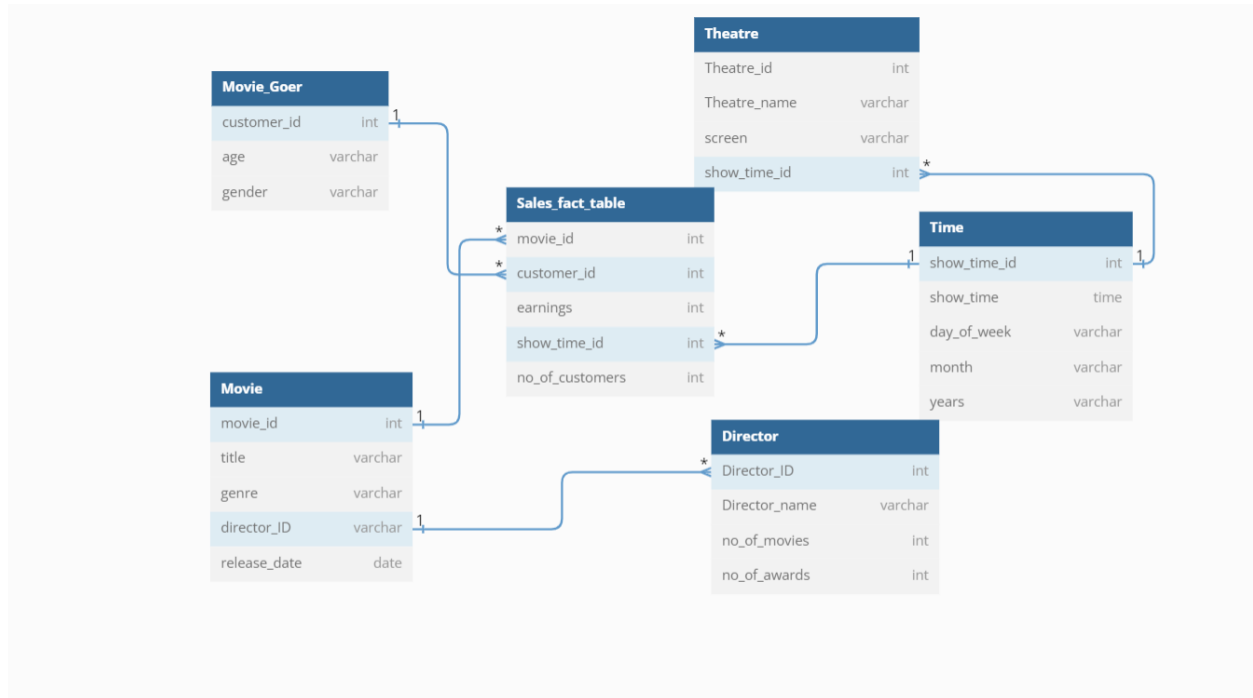
**Q2 B:**

```
SELECT m.title, sum(sf.revenue) as "Total Revenue"
from Sales_fact_table sf, time t, movie m
WHERE sf.show_time_id=t.show_time_id
and m.movie_id=sf.movie_id
and t.years = 2020
group by m.title;
```

**Q2 C:**

```
SELECT m.title, SUM(ts.ticket_price) AS "Total Revenue"
FROM ticket_sales ts
JOIN Movie m ON ts.movie = m.title
WHERE ts.year = 2020
GROUP BY m.title
```

## Q2 D:



## Query 1:

```
SELECT Theatre_name, screen, SUM(Sales_fact_table.earnings) AS total_earnings
FROM Theatre
JOIN Sales_fact_table ON Time.show_time_id = Sales_fact_table.show_time_id
JOIN Theatre ON time.show_time_id = Theatre.show_time_id
GROUP BY Theatre_name, screen;
```

## Query 2:

```
SELECT Director_name, no_of_movies
FROM Director
WHERE no_of_movies =
(SELECT MAX(no_of_movies)
FROM Director)
```

where Director\_name = "Christopher Nolan");

**Q3 A:**

1. Columnar Databases can easily handle data warehousing applications that involve the analysis of very large datasets. For example any retail industry produces a large volume of data such as sales done, customers information, transactions, inventory management, purchase orders etc. Columnar databases can be used to analyze all this data and identify hidden patterns and trends in customer behavior which could be beneficial for the business owners.
2. Instead of iterating over each row, the columnar database can read and process only specific columns needed for the query. This results in faster execution of the queries.
3. Columnar Database can store different type data such as structured, semi structured and unstructured data. For example social media sites generate a vast amount of unstructured data as well as semi structured data. The columnar database can combine all types of data and then perform the analysis to get the sentiment analysis on the post to understand the opinions of the users.
4. It provides various options to select the data source from such as different databases, file systems etc. It can be useful when the company having large data uses a particular type of data storage cloud. It gives them the liberty to use the data source accordingly.
5. Columnar database is a very good option for real time data processing and analysis as it can handle high speed transaction processing quite efficiently.

**Q3 B:**

1. IBM Cognos: It is BI tool providing different OLAP capabilities used for reporting, analytics and monitoring of events or metrics. It has different modules such as Cognos Planning, Cognos Controller and Cognos Analytics which can be used to retrieve information in different cases.
2. Oracle OLAP is an OLAP which provides various services for the Oracle Database. Custom functions can be made which comprises many small analytical functions which can be used to solve complex analytical problems.
3. Tableau : It is a data visualization tool handling large datasets easily. It is very user friendly and has simple drag and drop functions. It provides multidimensional as well as tabular models which can handle multiple dimensions and measures.
4. Mondrian: It is a java written OLAP engine. It uses MDX language for executing the queries. It also provides multidimensional analysis.
5. Microsoft SSAS: Microsoft SQL Server Analytic Services is a tool used for analysis and data mining to get information from multiple data sources. It also provides time-series analysis, regression analysis and clustering which can be useful while performing ad-hoc queries.

